



PgBouncer Deep Dive

Karen Jex | Senior Solutions Architect | Crunchy Data

June 2023

Agenda

- **What is PgBouncer?**
- **Connection Pooling Overview**
- **Getting Started with PgBouncer**
- **PgBouncer in Crunchy Data Products**
- **More PgBouncer Options**

What is PgBouncer?

- What does PgBouncer do?
- Why should I use it?

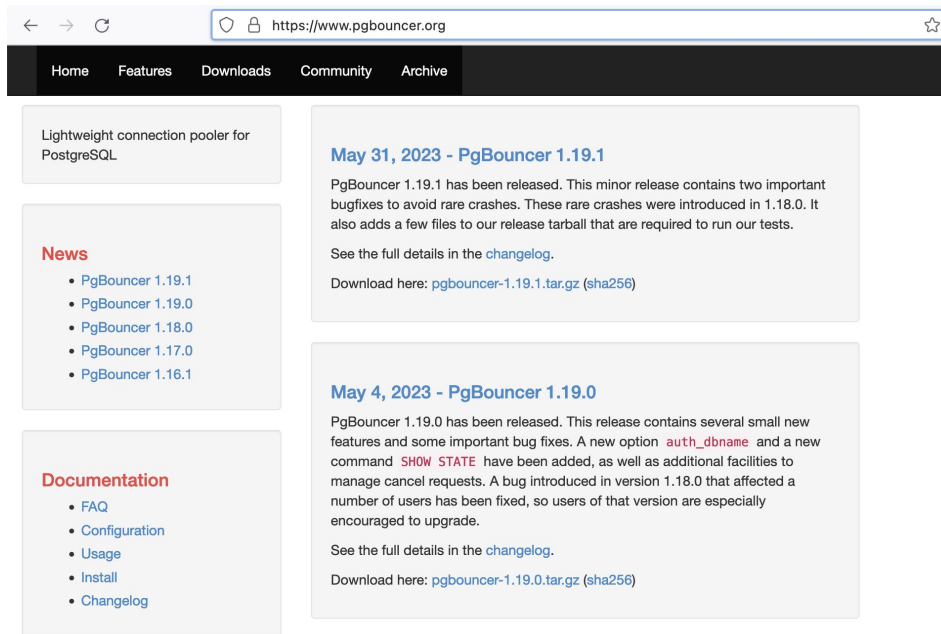
What is PgBouncer: **What does PgBouncer do?**

“lightweight, easy to configure connection pooler”



Image par Christine Sponchia de Pixabay

What is PgBouncer: PgBouncer Docs



<https://www.pgbouncer.org/usage.html#quick-start>

What is PgBouncer: Why should I use PgBouncer?



Crunchy PostgreSQL for Kubernetes

Cloud native containerized PostgreSQL with high availability

[v4 Docs](#) [v5 Docs](#) [Quick Start](#) [Product](#)
[GitHub](#) [Learn](#) [Video](#) [Blogs](#)

<https://access.crunchydata.com/documentation/pgbouncer/latest/>

Trusted Tools

[Crunchy PostgreSQL Operator](#) [pgBackrest](#)
[Crunchy Container Suite](#) [Patroni](#) [PostGIS](#)
[pgRouting](#) [pgAdmin4](#) [PgBouncer](#) [pgMonitor](#)



Crunchy HA PostgreSQL

Scripted Solutions for "Always On" data requirements

[Docs](#) [Product](#) [Video](#) [Blogs](#)

Trusted Tools

[PostgreSQL](#) [PgBouncer](#) [pgBackrest](#) [Patroni](#)
[psycopg2](#)



Crunchy Certified PostgreSQL

Crunchy Certified PostgreSQL is a trusted open-source relational database optimized for the enterprise. Certified for deployment on your choice of Platform, Infrastructure, or Cloud


[Docs](#) [Product](#) [PostgreSQL STIG](#)
[Videos](#) [Tutorial](#)

Trusted Tools

[PgBouncer](#) [pgBackrest](#) [pgMonitor](#)


Connection Pooling Overview

Connection Pooling Overview



max_connections = 100

Connection Pooling Overview



Connection Pooling Overview:

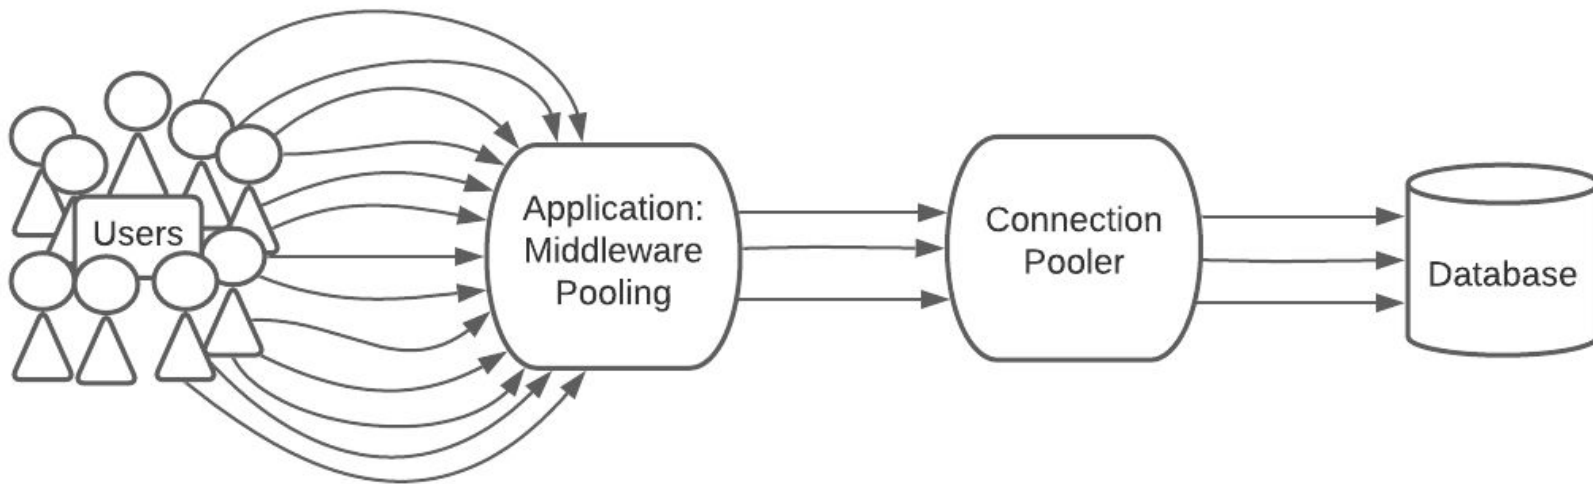
When not to use Connection Pooling

Consider potential disadvantages and whether or not it is required

- Additional latency
- Additional complexity
- Additional resources
- Security implications

Connection Pooling Overview:


When not to use Connection Pooling



Getting Started with PgBouncer

- Installing
- Configuring
- Monitoring

Getting Started with PgBouncer: Installing PgBouncer



<http://www.pgбouncer.org/faq.html#should-pgбouncer-be-installed-on-the-web-server-or-database-server>

Getting Started with PgBouncer: **Installing PgBouncer**

<https://www.pgbouncer.org/downloads/>

RPM: yum.postgresql.org

Deb: apt.postgresql.org

<https://github.com/pgbouncer/pgbouncer.git>

Getting Started with PgBouncer: Installing PgBouncer

```
[karen_jex@rocky9 ~]$ sudo dnf -y install pgbouncer
```

```
...
```

```
=====
Package            Architecture  Version           Repository        Size
=====
```

```
Installing:
```

```
pgbouncer          x86_64        1.19.1-42.rhel9   pgdg-common       215 k
```

```
Installing dependencies:
```

```
python3-psycopg2   x86_64        2.9.5-1.rhel9     pgdg-common       188 k
```

```
...
```

```
Installed:
```

```
pgbouncer-1.19.1-42.rhel9.x86_64    python3-psycopg2-2.9.5-1.rhel9.x86_64
```

```
Complete!
```

Getting Started with PgBouncer: **Configuring PgBouncer**

```
$ cat /etc/pgbouncer/pgbouncer.ini
;;;
;;; PgBouncer configuration file
;;;

;; database name = connect string
;;
;; connect string params:
;;   dbname= host= port= user= password= auth_user= client_encoding= datestyle= timezone=
;;   pool_size= reserve_pool= max_db_connections= pool_mode= connect_query= application_name=
...
;; Read additional config from other file
;%include /etc/pgbouncer/pgbouncer-other.ini
```


Getting Started with PgBouncer: **Configuring PgBouncer**

databases section

```
[databases]

charleston = host=localhost dbname=charleston

;charleston = dbname=charleston host=192.168.152.128 port=5433

;ch = dbname=charleston host=localhost
```

Getting Started with PgBouncer: **Configuring PgBouncer**

Authentication settings

```
;;;
;;; Authentication settings
;;;

;; any, trust, plain, md5, cert, hba, pam
auth_type = scram-sha-256

;auth_hba_file =
```

Getting Started with PgBouncer: **Configuring PgBouncer**

Authentication settings

PgBouncer handles its own client authentication and **has its own database of users**.

```
auth_file = /etc/pgbouncer/userlist.txt
```

```
$ echo '"karen" "mysupersecretpassword"' > /etc/pgbouncer/userlist.txt
```

Getting Started with PgBouncer: **Configuring PgBouncer**

Authentication settings

```
[databases]
```

```
charleston = host=localhost dbname=Charleston auth_user=karen auth_dbname=authdb
```

```
[pgbouncer]
```

```
;; Authentication database that can be set globally to run "auth_query".
```

```
auth_dbname = authdb
```

```
;; Query to use to fetch username and password hash from database.
```

```
auth_query = SELECT username, passwd FROM pg_shadow WHERE username=$1
```

Getting Started with PgBouncer: **Configuring PgBouncer**

Admin User

```
;;;
;;; Users allowed into database 'pgbouncer'
;;;

;; comma-separated list of users who are allowed to change settings
admin_users = postgres
admin_users = karen
```

Getting Started with PgBouncer: **Configuring PgBouncer**

Type of Pooling

When can a server connection be returned to the pool for reuse?

```
pool_mode = session | transaction | statement
```

Getting Started with PgBouncer: **Configuring PgBouncer**

pool_mode = session (default)

The connection is released back to the pool after the client disconnects

Getting Started with PgBouncer: **Configuring PgBouncer**

pool_mode = transaction

The connection is released back to the pool after the transaction finishes

Getting Started with PgBouncer: **Configuring PgBouncer**

pool_mode = statement

The connection is released back to the pool after the query finishes

Transactions spanning multiple statements not allowed

Tends to break applications, especially if they've not been designed with this in mind

Getting Started with PgBouncer: **Configuring PgBouncer**

Pool Sizing

Parameter	Description	Default value
max_client_conn	Max number of client connections allowed	100
default_pool_size	Number of server connections per user/database pair	20
min_pool_size	Add server connections to pool if below this number	0 (disabled)
reserve_pool_size	How many additional connections to allow to a pool	0 (disabled)
reserve_pool_timeout	Use connection from the reserve pool if client waits this long (s)	5.0
max_db_connections	Max number of connections allowed per database	0 (unlimited)
max_user_connections	Max number of connections allowed per user	0 (unlimited)


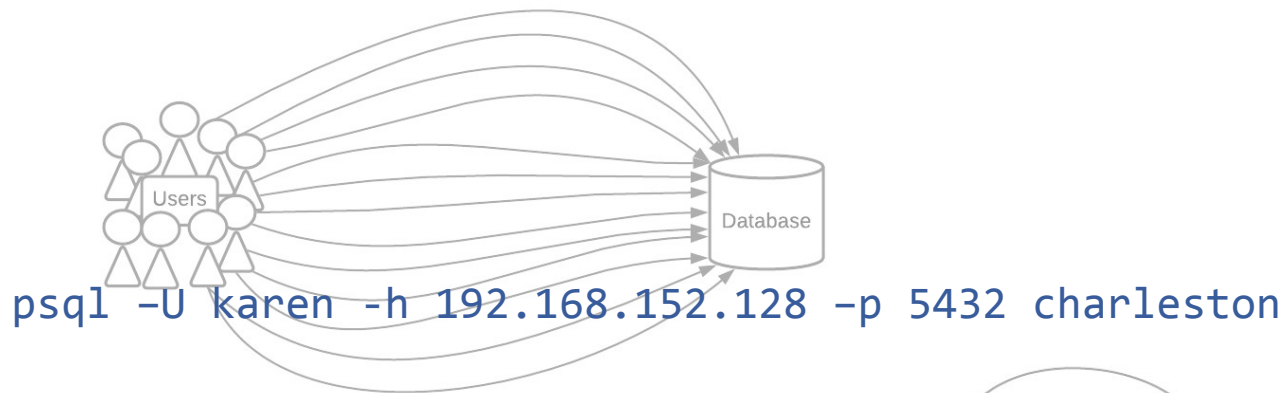
Getting Started with PgBouncer: Starting PgBouncer

```
$ sudo systemctl start pgbouncer
```

or

```
$ pgbouncer -d /etc/pgbouncer/pgbouncer.ini
```

Getting Started with PgBouncer: Connecting via PgBouncer



Getting Started with PgBouncer: Connecting via PgBouncer

```
$ psql -U [username] -h [pgbouncer_host] -p [pgbouncer_port] [pgbouncer_db]
```

```
$ psql -U karen -h 192.168.152.128 -p 6432 charleston
```

PgBouncer Console

```
$ psql -U postgres -h 127.0.0.1 -p 6432 pgbouncer
```

Getting Started with PgBouncer: **Monitoring**

```
pgbouncer=# SHOW HELP;
```

```
SHOW HELP|CONFIG|DATABASES|POOLS|CLIENTS|SERVERS|USERS|VERSION
```

```
SHOW PEERS|PEER_POOLS
```

```
SHOW FDS|SOCKETS|ACTIVE_SOCKETS|LISTS|MEM|STATE
```

```
SHOW DNS_HOSTS|DNS_ZONES
```

```
SHOW STATS|STATS_TOTALS|STATS_AVERAGES|TOTALS
```

```
...
```

Getting Started with PgBouncer: Monitoring

```
pgbouncer=# SHOW clients;
```

type	user	database	state	port	connect_time	application_name
-----+-----+-----+-----+-----+-----+-----						
C	karen	charleston	active	6432	2023-06-12 15:15:05 CEST	psql
C	karen	pgbouncer	active	6432	2023-06-12 15:14:06 CEST	psql

Getting Started with PgBouncer: **pgbouncer_fdw**

“pgbouncer_fdw provides a direct SQL interface to the PgBouncer SHOW commands. It takes advantage of the dblink_fdw feature to provide a more typical, table-like interface to the current status of your PgBouncer server(s). This makes it easier to set up monitoring or other services that require direct access to PgBouncer statistics.”

<https://www.crunchydata.com/blog/making-pgbouncer-easier-to-monitor>

https://github.com/CrunchyData/pgbouncer_fdw

Getting Started with PgBouncer: Admin

```
pgbouncer=# SHOW HELP;
```

```
...
```

```
RELOAD
```

```
PAUSE [<db>]
```

```
RESUME [<db>]
```

```
DISABLE <db>
```

```
ENABLE <db>
```

```
RECONNECT [<db>]
```

```
KILL <db>
```

```
SUSPEND
```

```
SHUTDOWN
```

```
WAIT_CLOSE [<db>]
```

PgBouncer in Crunchy Data Products

- **Crunchy Postgres**

full service scripted solution for deploying Production PostgreSQL
anywhere you want to run it

- **Crunchy Postgres for Kubernetes**

Cloud Native Postgres on Kubernetes
powered by Crunchy Postgres Operator (PGO)

- **Crunchy Bridge**

Fully managed Postgres on your choice of Cloud provider

PgBouncer in Crunchy Data Products: **Crunchy Postgres**

Implemented via the pgbouncer role (extracts from docs)

databases

collect_pgbouncer_metrics

VARIABLE NAME:	collect_pgbouncer_metrics
DESCRIPTION:	Whether to collect PgBouncer metrics for pgMonitor
DEFAULT VALUE:	false
TYPE:	boolean
REQUIRED:	false
EXAMPLE VALUE(S):	N/A
NOTES:	When <code>true</code> , the <code>pgbouncer</code> role will assume that the 'dblink' and 'pgbouncer_fdw' extensions are installed in the target database (the <code>pgmonitor</code> role will do this if <code>collect_pgbouncer_metrics: true</code> is set there) and create the Foreign Data Wrapper server and user mappings for the <code>pgmonitor_db_user</code> (usually <code>ccp_monitoring</code>). This presents a set of views in the PostgreSQL database that pgMonitor will use to collect metric data for monitoring PgBouncer itself.

VARIABLE NAME:	databases
DESCRIPTION:	Configured user databases for PgBouncer
DEFAULT VALUE:	[]
TYPE:	list of dictionaries
REQUIRED:	false
EXAMPLE VALUE(S):	► Complex value; click to see in full

PgBouncer in Crunchy Data Products:

Crunchy Postgres for Kubernetes

<https://access.crunchydata.com/documentation/postgres-operator/latest/quickstart/#connect-to-the-postgres-cluster>

<https://access.crunchydata.com/documentation/postgres-operator/latest/tutorial/connection-pooling/>


```
spec:
  proxy:
    pgBouncer:
      config:
        global:
          pool_mode: transaction
```

More PgBouncer Options

- Multiple Instances
- Read-Write / Read Only Routing
- Load Balancing
- Fault Tolerance
- Taming "Badly Behaved" Applications

More PgBouncer Options: **Multiple Instances**

SO_REUSEPORT



More PgBouncer Options: **Multiple Instances**

- Create multiple PgBouncer instances
- Use SO_REUSEPORT so multiple PgBouncer instances listen on the same port
- Use more than one thread
- easy scaling!
- so_reuseport info:
<https://www.pgбouncer.org/config.html#low-level-network-settings>

More PgBouncer Options: **Multiple Instances**

```
[pgbouncer]  
...  
so_reuseport = 1
```

```
$ cp /etc/pgbouncer.service /etc/systemd/system/pgbouncer@.service  
  
$ daemon reload  
  
$ systemctl start pgbouncer@1  
$ systemctl start pgbouncer@2  
$ etc
```



More PgBouncer Options: **Read-Write / Read Only Routing**

Use different pgBouncer databases to:

- Route read-write traffic to the current primary database
- Route read-only traffic to a standby database

```
[pgbouncer]
...
charleston_rw = dbname=charleston host=<virtual_IP>
charleston_ro = dbname=charleston host=<standby_ip>
```

More PgBouncer Options: Read-Write / Read Only Routing



More PgBouncer Options: **Read-Write / Read Only Routing**

```
pgbouncer1.ini
```

```
[databases]
```

```
read_write = dbname=charleston host=<Virtual IP>
```

```
read_only = dbname=charleston host=192.168.152.129
```


```
pgbouncer2.ini
```

```
[databases]
```

```
read_write = dbname=charleston host=<Virtual IP>
```

```
read_only = dbname=charleston host=192.168.152.130
```

More PgBouncer Options: Load Balancing



More PgBouncer Options: Fault Tolerance

When using `SO_REUSEPORT`

- Connections to databases created on a round-robin basis
- Database connections are persistent

If PgBouncer is on a **database server**

If DB server/PgBouncer instance is lost:

- Connection between application and database terminated
- Round-robin begins again

If PgBouncer on an **external** machine

- if the database goes down:
 - connection to PgBouncer remains
 - PgBouncer attempts to connect to another endpoint
 - the application connection is not disrupted

More PgBouncer Options: Taming “Badly Behaved” Applications

- Database traffic from certain application may be unpredictable
- Use PgBouncer to tame “badly behaved” applications by using distinct port numbers
- Provide database connections to deal with sudden bursts of incoming traffic in different ways
- prevent the database from becoming swamped during high activity periods

More PgBouncer Options: Taming “Badly Behaved” Applications

```
%include /etc/pgbouncer/pgbouncer.ini

[databases]

toomuch = dbname=pgbench host=127.0.0.1

[users]

bounce = pool_mode=transaction max_user_connections=200

[pgbouncer]

listen_port = 6433

max_db_connections = 10
```

More PgBouncer Options: Taming “Badly Behaved” Applications

```
$ /usr/bin/pgbench -h localhost -p 6433 -U bounce -C -c 200 -T 600 toomuch
```

```
connection to database "toomuch" failed:
```

```
ERROR: no more connections allowed (max_client_conn)
```

```
client 100 aborted while establishing connection
```

```
connection to database "toomuch" failed:
```

```
ERROR: no more connections allowed (max_client_conn)
```

```
client 100 aborted while establishing connection
```

```
connection to database "toomuch" failed:
```

```
ERROR: no more connections allowed (max_client_conn)
```

```
client 100 aborted while establishing connection
```




That's All Folks!